


Analysis of the Training and Validation Losses:



Loss and Val_loss:

Training Loss:

The model's training loss started at 0.2155 in the first epoch and gradually decreased to 0.0015 by the 50th epoch. This consistent decrease indicates that the model was learning and improving its ability to fit the training data.

Validation Loss:

The validation loss showed a similar trend, starting at 0.5540 in the first epoch and decreasing to 0.0022 by the 50th epoch. The validation loss follows the training loss closely, indicating that the model is generalizing well to the validation data without overfitting significantly.

Learning Rate Adjustments:

Initially, the learning rate was set at 0.001. As training progressed, the learning rate was halved to 0.0005 from epoch 44 onwards. This reduction in learning rate likely helped in fine-tuning the model and stabilizing the loss, as reflected in the stable and low validation loss in the final epochs.

as reflected in the stable and low validation loss in the final epochs.

Convergence:

The model's loss values stabilized towards the later epochs, particularly after epoch 40, where both training and validation losses were consistently low. This suggests that the model had converged and reached an optimal point in the training process.

Evaluation Metrics:

RMSE (Root Mean Square Error): 29.57

RMSE gives an indication of the average magnitude of the errors in the predictions. Lower RMSE values are better, indicating that the model's predictions are close to the actual values. An RMSE of 29.57 is reasonably low, suggesting good predictive accuracy.

MAE (Mean Absolute Error): 380.44

MAE is the average of the absolute differences between predicted and actual values. While MAE is typically lower than RMSE, it gives a direct interpretation of average error. Here, an MAE of 380.44 suggests that, on average, predictions deviate by about 380.44 units from the actual values.

R-squared (R^2): 0.891

The R^2 value of 0.891 indicates that 89.1% of the variance in the target variable is explained by the model. This is a strong R^2 value, indicating that the model fits the data well.

MAPE (Mean Absolute Percentage Error): 2.94%

MAPE indicates the average percentage error between the predicted and actual values. A MAPE of 2.94% is quite low, signifying that, on average, the model's predictions are off by less than 3% relative to the actual values.

Summary:

Model Performance: The model demonstrates strong performance with low training and validation losses, a high R^2 value, and low RMSE, MAE, and MAPE values. These metrics indicate that the model is both accurate and generalizes well to unseen data.

Learning Dynamics: The steady decrease in loss values and the appropriate adjustment of the learning rate indicate that the training process was well-tuned, leading to effective learning without significant overfitting.

Final Model State: The final model, based on these results, appears to be well-trained and capable of making accurate predictions, with minimal errors and a high degree of variance explained by the model.

